

Tiny Transformer with MultiHeadAttention

We propose MultiHeadAttention using scaled dot-product attention.

TinyTransformerBlock stacks attention and a feed-forward network.

LayerNorm and GELU stabilize TinyTransformerBlock training.

The train_step function optimizes mean squared loss on batch tensors.